



CASTER-ZT Component Training Convergence


(a) Trust Autoencoder (AE_ϕ , $\sim 2.5K$ params)



(b) Risk Scorer (R_ϕ , $\sim 1.5K$ params)



(c) Contrastive Safety Encoder (q_ϕ , $\sim 8K$ params)



(d) GNN Encoder + Policy (f_θ, π_θ , $\sim 37K$ params)

